

## **INDEX**

<b>S.No.</b>	<b>DESCRIPTION</b>	<b>Page No.</b>
1.	Institute Vision and Mission	2
2.	Department Vision and Mission	3
3.	PEO's	4
4.	PO's, PSO's	5
5.	List of CO with BTL and mapping with PO/ PSO	7
6.	CO/ PO/ PSO Mapping Justification	10
7.	Syllabus	11
8.	List of Experiments – mappingwith CO/PO/PSO	12
9.	Introduction to the lab	14
<b>EXPERIMENTS</b>		
1.	Implementation of Tokenization using UNIX commands	15
2.	Implement a word tokenization using regular expressions	18
3.	Implement Minimum Edit Distance(MED) algorithm for spelling correction	20
4.	Implement n-gram language model	22
5.	Implement Naïve Bayes classification for sentiment analysis	24
6.	Implement POS tagging using HMM	26
7.	Implementation of Implement CKY parsing algorithm	28
8.	Implement PCKY parsing algorithm	30
9.	Implementation of Computing cosine similarity between the words using term document matrix and term-term matrix	34
10.	Implementation of TF-IDF matrix for the given document set	36
11.	Implementation of Language Model Using Feed forward Neural Network	38
12.	Implement Language Model Using RNN	40
13.	Implement Perform Text Analytics	43

**VASAVI COLLEGE OF ENGINEERING (AUTONOMOUS)**  
**IBRAHIMBAGH, HYDERABAD – 500031**

**Institute Vision**

Striving for a symbiosis of technological excellence and human values.

**Institute Mission**

To arm young brains with competitive technology and nurture holistic development  
of the individuals for a better tomorrow.

**VASAVI COLLEGE OF ENGINEERING (AUTONOMOUS)**  
**IBRAHIMBAGH, HYDERABAD – 500031**

**DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING**

**Department Vision**

To be a center for academic excellence in the field of Computer Science and Engineering education to enable graduates to be ethical and competent professionals.

**Department Mission**

To enable students to develop logic and problem solving approach that will help build their careers in the innovative field of computing and provide creative solutions for the benefit of society.

**VASAVI COLLEGE OF ENGINEERING (AUTONOMOUS)**  
**IBRAHIMBAGH, HYDERABAD – 500031**

**DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING**

**PEO's, PO's and PSO's**

**PROGRAM EDUCATIONAL OBJECTIVES (PEOs):**

Graduates should be able to utilize the knowledge gained from their academic program to:

**PEO 1:** Solve problems in a modern technological society as valuable and productive engineers.

**PEO 2:** Function and communicate effectively, both individually and within multidisciplinary teams.

**PEO 3:** Be sensitive to the consequences of their work, both ethically and professionally, for productive professional careers.

**PEO 4:** Continue the process of life-long learning.

**PROGRAM OUTCOMES (POs):**

Engineering Graduates will be able to:

**PO1: Engineering knowledge:** Apply the knowledge of mathematics, science, engineering fundamentals, and an engineering specialization to the solution of complex engineering problems.

**PO2: Problem analysis:** Identify, formulate, review research literature, and analyze complex engineering problems reaching substantiated conclusions using first principles of mathematics, natural sciences, and engineering sciences.

**PO3: Design/development of solutions:** Design solutions for complex engineering problems and design system components or processes that meet the specified needs with appropriate

consideration for the public health and safety, and the cultural, societal, and environmental considerations.

**PO4: Conduct investigations of complex problems:** Use research-based knowledge and research methods including design of experiments, analysis and interpretation of data, and synthesis of the information to provide valid conclusions.

**PO5: Modern tool usage:** Create, select, and apply appropriate techniques, resources, and modern engineering and IT tools including prediction and modeling to complex engineering activities with an understanding of the limitations.

**PO6: The engineer and society:** Apply reasoning informed by the contextual knowledge to assess societal, health, safety, legal and cultural issues and the consequent responsibilities relevant to the professional engineering practice.

**PO7: Environment and sustainability:** Understand the impact of the professional engineering solutions in societal and environmental contexts, and demonstrate the knowledge of, and need for sustainable development.

**PO8: Ethics:** Apply ethical principles and commit to professional ethics and responsibilities and norms of the engineering practice.

**PO9: Individual and team work:** Function effectively as an individual, and as a member or leader in diverse teams, and in multidisciplinary settings.

**P10: Communication:** Communicate effectively on complex engineering activities with the engineering community and with society at large, such as, being able to comprehend and write effective reports and design documentation, make effective presentations, and give and receive clear instructions.

**P11: Project management and finance:** Demonstrate knowledge and understanding of the engineering and management principles and apply these to one's own work, as a member and leader in a team, to manage projects and in multidisciplinary environments.

**P12: Life-long learning:** Recognize the need for, and have the preparation and ability to engage in independent and life-long learning in the broadest context of technological change.

**PROGRAM SPECIFIC OUTCOMES (PSOs):**

**PSO1:** Graduates will have knowledge of programming and designing to develop solutions for engineering problems.

**PSO2:** Graduates will be able to demonstrate an understanding of system architecture, information management and networking.

**PSO3:** Graduates will possess knowledge of applied areas of computer science and engineering and execute them appropriately.



**VASAVI COLLEGE OF ENGINEERING (Autonomous)**  
**IBRAHIMBAGH, HYDERABAD – 500 031**

**Department of Computer Science & Engineering**

**UI21PE751CS Natural Language Processing Lab**

**B.E VII SEMESTER**

**Course Outcomes mapping with Blooms Taxonomy**

S.NO	Course Outcomes	BTL
1	Implement the word tokenization for NLP pre-processing	3
2	Implement the N-gram and Neural language models	3
3	Implement parsing algorithms for grammar checking	3
4	Implement solution for sentiment analysis	3
5	Implement the NLP task namely document classification	3



**VASAVI COLLEGE OF ENGINEERING (Autonomous)**  
**IBRAHIMBAGH, HYDERABAD – 500 031**

**Department of Computer Science & Engineering**

**UI21PE751CS Natural Language Processing Lab**

**B.E VII SEMESTER**

**CO- PO Mapping**

<b>Course Code</b>	<b>Course Outcomes</b>	<b>PO1</b>	<b>PO2</b>	<b>PO3</b>	<b>PO4</b>	<b>PO5</b>	<b>PO6</b>	<b>PO7</b>	<b>PO8</b>	<b>PO9</b>	<b>PO1 0</b>	<b>PO1 1</b>	<b>PO1 2</b>
CO1	Implement the word tokenization for NLP pre-processing				3	3							
CO2	Implement the N-gram and Neural language models				3	3							
CO3	Implement parsing algorithms for grammar checking				3	3							
CO4	Implement solution for sentiment analysis				3	3							
CO5	Implement the NLP task namely document classification				3	3							
	Average				3	3							



**VASAVI COLLEGE OF ENGINEERING (Autonomous)**  
**IBRAHIMBAGH, HYDERABAD – 500 031**

**Department of Computer Science & Engineering**

**UI21PE751CS Natural Language Processing Lab**

**B.E VII SEMESTER**

**CO-PSO Mapping**

<b>Course Code</b>	<b>Course Outcomes</b>	<b>PSO1</b>	<b>PSO2</b>	<b>PSO3</b>
CO1	Implement the word tokenization for NLP pre-processing			3
CO2	Implement the N-gram and Neural language models			3
CO3	Implement parsing algorithms for grammar checking			3
CO4	Implement solution for sentiment analysis			3
CO5	Implement the NLP task namely document classification			3
	Average			3



**VASAVI COLLEGE OF ENGINEERING (Autonomous)**  
**IBRAHIMBAGH, HYDERABAD – 500 031**

**Department of Computer Science & Engineering**

**UI21PE751CS Natural Language Processing Lab**

**B.E VII SEMESTER**

**CO-PO Mapping Justification**

UI21PE751CS	CO1	PO4	3	PO4 expects the students to have research-based knowledge and use research methods for analysing and interpreting data and to provide valid conclusions.
		PO5	3	PO5 expects the students to be able to predict and model complex engineering activities with an understanding of the limitations
	CO2	PO4	3	PO4 expects the students to have research-based knowledge and use research methods for analysing and interpreting data and to provide valid conclusions.
		PO5	3	PO5 expects the students to be able to predict and model complex engineering activities with an understanding of the limitations
	CO3	PO4	3	PO4 expects the students to have research-based knowledge and use research methods for analysing and interpreting data and to provide valid conclusions.
		PO5	3	PO5 expects the students to be able to predict and model complex engineering activities with an understanding of the limitations
	CO4	PO4	3	PO4 expects the students to have research-based knowledge and use research methods for analysing and interpreting data and to provide valid conclusions.
		PO5	3	PO5 expects the students to be able to predict and model complex engineering activities with an understanding of the limitations
	CO5	PO4	3	PO4 expects the students to have research-based knowledge and use research methods for analysing and interpreting data and to provide valid conclusions.
		PO5	3	PO5 expects the students to be able to predict and model complex engineering activities with an understanding of the limitations

**CO-PSO Mapping Justification**

UI21PE751CS	CO1	PSO3	3	PSO3 expects the students to be able to gain the knowledge of computer science and engineering in the areas of Data Analytics and apply them.
	CO2	PSO1	3	PSO3 expects the students to be able to gain the knowledge of computer science and engineering in the areas of Data Analytics and apply them.
	CO3	PSO1	3	PSO3 expects the students to be able to gain the knowledge of computer science and engineering in the areas of Data Analytics and apply them.
	CO4	PSO1	3	PSO3 expects the students to be able to gain the knowledge of computer science and engineering in the areas of Data Analytics and apply them.
	CO5	PSO1	3	PSO3 expects the students to be able to gain the knowledge of computer science and engineering in the areas of Data Analytics and apply them.



**VASAVI COLLEGE OF ENGINEERING (Autonomous)**  
**IBRAHIMBAGH, HYDERABAD – 500 031**

**Department of Computer Science & Engineering**

**UI21PE751CS NATURAL LANGUAGE PROCESSING LAB**

**SYLLABUS FOR B.E VII SEMESTER**

<b>Course Objectives</b>	<b>Course Outcomes</b>
Students should be able to	At the end of the course, students will be able to
<ul style="list-style-type: none"> <li>1. Apply dynamic programming design strategy for NLP tasks .</li> <li>2. Design Feedforward and Recurrent Neural Networks for performing Language Modelling</li> </ul>	<ul style="list-style-type: none"> <li>1. Implement the word tokenization for NLP pre-processing</li> <li>2. Implement the N-gram and Neural language models.</li> <li>3. Implement parsing algorithms for grammar checking</li> <li>4. Implement solution for sentiment analysis</li> <li>5. Implement the NLP task namely document classification</li> </ul>

1. Implementation of Tokenization using UNIX commands.
2. Implement a word tokenization using regular expressions.
3. Implement Minimum Edit Distance(MED) algorithm for spelling correction.
4. Implement n-gram language model.
5. Implement Naïve Bayes classification for sentiment analysis.
6. Implement POS tagging using HMM.
7. Implementation of Implement CKY parsing algorithm.
8. Implement PCKY parsing algorithm.
9. Implementation of Computing cosine similarity between the words using term document matrix and term-term matrix.
10. Implementation of TF-IDF matrix for the given document set.
11. Implementation of Language Model Using Feed forward Neural Network.
12. Implement Language Model Using RNN.
13. Perform Text Analytics



**VASAVI COLLEGE OF ENGINEERING (Autonomous)**  
**IBRAHIMBAGH, HYDERABAD – 500 031**

**Department of Computer Science & Engineering**

**UI21PE751CS Natural Language Processing Lab**

**B.E VII SEMESTER**

**Experiments to CO, PO, PSO Mapping**

S. No.	Name of the Experiment	CO	PO	PSO
1	<b>Implementation of Tokenization using UNIX commands</b>	1	4,5	3
2	<b>Implement a word tokenization using regular expressions</b>	1	4,5	3
3	<b>Implement Minimum Edit Distance(MED) algorithm for spelling correction</b>	1	4,5	3
4	<b>Implement n-gram language model</b>	1	4,5	3
5	<b>Implement Naïve Bayes classification for sentiment analysis</b>	3	4,5	3
6	<b>Implement POS tagging using HMM</b>	2	4,5	3
7	<b>Implementation of Implement CKY parsing algorithm</b>	5	4,5	3
8	<b>Implement PCKY parsing algorithm</b>	5	4,5	3
9	<b>Implementation of Computing cosine similarity between the words using term document matrix and term-term matrix</b>	2	4,5	3
10	<b>Implementation of TF-IDF matrix for the given document set</b>	2	4,5	3
11	<b>Implementation of Language Model Using Feed forward Neural Network</b>	4	4,5	3
12	<b>Implement Language Model Using RNN</b>	4	4,5	3

<b>13</b>	<b>Implement Perform Text Analytics</b>	<b>3</b>	<b>4,5</b>	<b>3</b>
-----------	---	----------	------------	----------



**VASAVI COLLEGE OF ENGINEERING (Autonomous)  
IBRAHIMBAGH, HYDERABAD – 500 031**

**Department of Computer Science & Engineering**

**UI21PE751CS Natural Language Processing Lab**

**B.E VII SEMESTER**

### **Introduction to the lab**

Natural Language Processing (NLP) is a field of artificial intelligence concerned with the interaction between computers and human language. It involves the development of algorithms and models that enable computers to understand, interpret, and generate human language in a way that is meaningful. NLP encompasses a wide range of tasks and applications, including:

- 1. Text Classification and Categorization:** Assigning labels or categories to text based on its content (e.g., sentiment analysis, topic classification).
- 2. Named Entity Recognition (NER):** Identifying and categorizing entities (such as names of people, organizations, locations) in text.
- 3. Text Generation:** Creating coherent and contextually relevant text based on given input (e.g., chatbots, language translation).
- 4. Machine Translation:** Translating text from one language to another automatically (e.g., Google Translate).
- 5. Sentiment Analysis:** Determining the sentiment expressed in a piece of text (positive, negative, neutral).
- 6. Speech Recognition:** Converting spoken language into text (e.g., virtual assistants like Siri or Alexa).
- 7. Information Retrieval:** Finding relevant information in a large collection of text (e.g., search engines).

## EXPERIMENT-1 : Implementation of Tokenization using UNIX commands

Certainly! Word tokenization is a fundamental task in Natural Language Processing (NLP) that involves breaking down a text into individual words or tokens. We can create a simple program using Unix utilities like sed, tr, grep, and awk to perform basic word tokenization. Here's a step-by-step example of how you can achieve

1. Create a text file text1.txt containing text of  
your choice  
Ans:echo  
"Thisisanexampletextfile.">text1.txt

2. Use tr utility to replace every occurrence of a given character to another  
given character  
Ans:tr'o"x'<example.txt(HelloWorld!) >HelloxWxrld!

3. What are s and c options of tr  
command?  
Ans:echo "aaabbbccc" |tr-  
s'a'Output:abbcc

**-s option (squeeze-repeats):** This option is used to squeeze (reduce) repeated characters in the input to a single character.

**-c option (complement):** This option is used to complement the set of characters. When used, it replaces characters not listed in the first set with the corresponding character from the second set

echo"abc123" |tr -c'a-z"X"-->Output:XXX123

4. What is the output of the following command? tr-a-zA-Z<text1.txt

Ans: The command tr-a-zA-Z<text1.txt will transform all lowercase letters to uppercase in the content of the file text1.txt. The output will be displayed on the terminal.

5. Create a text file text2.txt containing any text (one word per line) sort the words in text2 file using “sort” command

Ans:echo-e"banana\napple\norange\nkiwi\ngrape">text2.txt sort text2.txt

6. Sort and display unique lines in the  
text2.txt file  
Ans:sort -u text2.txt

Output:

apple

banana

grape

kiwi

orange

7. Sort and display unique lines in text2.txt file such that each word is preceded with frequency count of the word in text2.txt file

Ans: sort text2.txt | uniq -c

Output:

1 apple

2 banana

1 grape

3 kiwi

1 orange

8. Obtain and display the tokens in text1.txt file

Ans: tr -sc 'A-Za-z' '\n' < text1.txt | grep -v '^\$' (input: This is an example text file.)

Output:

This is an

example text

file

9. Display the tokens in sorted order

Ans: tr -sc 'A-Za-z' '\n' < text1.txt | grep -v '^\$' | sort

---

Output:

This an

example file

is text

10. Display the unique tokens in sorted order

Ans: tr -sc 'A-Za-z' '\n' < text1.txt | grep -v '^\$' | sort | uniq

Output:

This an

example file

is text

## EXPERIMENT-2 : Implement a word tokenization using regular expressions

To tokenize words using the regular expression tokenizer provided by NLTK (Natural Language Toolkit) in Python, you'll first need to ensure NLTK is installed (pip install nltk). Then, you can use the RegexpTokenizer class from NLTK to define and apply custom tokenization rules based on regular expressions. Here's how you can create a Python program to tokenize words using NLTK's regular expression tokenizer

```
from nltk.tokenize import RegexpTokenizer  
  
s = "Good muffins cost $3.88\nin New York. Please buy me\n\ttwo of  
them.\n\nThanks."  
tokenizer = RegexpTokenizer(r'\w+|\$[\d\.]+|\S+')  
  
print(tokenizer.tokenize(s))
```

**Output:**

```
[Good, muffins, cost, '$3.88', 'in', 'New', 'York', '.', 'Please', 'buy', 'me', 'two', 'of', 'them', '.', 'Thanks', '.']
```

```
import nltk  
nltk.download('punkt')  
nltk.download('averaged_perceptron_t  
agger')  
nltk.download('maxent_ne_chu  
nker')  
nltk.download('words')  
  
sentence="""At eight o'clock on Thursday morning... Arthur didn't  
feel very good."""  
tokens = nltk.word_tokenize(sentence)  
  
print(tokens)
```

**Output:**

```
['At', 'eight', 'o''clock', 'on', 'Thursday', 'morning', '...', 'Arthur', 'did', "n't", 'feel', 'very', 'good', '.']
```

```
text = "This is V.C.E C.S.E I am a Student . I paid a fees of 13.0"  
print(re.findall("(?:[A-Z]\. )+[A-Z]",text)) # ([A-Z]\. )+ add ":" before this for concatenation  
text = "That U.S.A poster-print costs $12.40... which is 3.45."  
print(re.split(" ",text)) #regex for hyphenated-words  
print(re.findall("\w+-\w+",text))  
print(re.findall("(?:\w+-\w+)",text))  
from nltk.corpus import stopwords  
print(word_tokenize(text))
```

```
li = []
for w in word_tokenize(text):
    if w not in stopwords.words('english'): li.append(w)
li
Output:
['V.C.E', 'C.S.E']
['That', ", 'U.S.A', 'poster-print', 'costs', '$12.40...', 'which', 'is', '3.45.',] ['poster-print']
['poster-print']
['That', 'U.S.A', 'poster-print', 'costs', '$', '12.40', '...', 'which', 'is', '3.45', '.']
['That', 'U.S.A', 'poster-print', 'costs', '$', '12.40', '...', '3.45', '.']
```

**EXPERIMENT - 3: Implement Minimum Edit Distance(MED) algorithm for spelling correction**

To compute the Minimum Edit Distance (also known as Levenshtein distance) using NLTK in Python, you can use the `edit_distance` function provided by NLTK's `edit_distance` module. This function calculates the minimum number of single-character edits (insertions, deletions, substitutions) required to change one word into another. Here's a Python program to demonstrate how to compute the Minimum Edit Distance using NLTK.

```
from re import M

#source = input('Enter the string given: ') source = 'kitten'
#target = input('Enter the string to convert to: ') target = 'sitting'

m = len(source) n = len(target)

dp = [[0 for i in range(m+1)] for j in range(n+1)] print(dp)

for i in range(m+1):
    dp[0][i] = i

for j in range(n+1):
    dp[j][0] = j print(dp)

for i in range(1, n+1):
    for j in range(1, m+1):
        if source[j-1] == target[i-1]:
            cost = 0 else:
            cost = 1

        dp[i][j] = min(dp[i-1][j]+1, dp[i][j-1]+1, dp[i-1][j-1]+cost) print(dp)

print('Edit Distance: ', dp[n][m])
```

**Output:**

```
[[0,0,0,0,0,0],[0,0,0,0,0,0],[0,0,0,0,0,0],
 [[0,1,2,3,4,5,6],[1,0,0,0,0,0],[2,0,0,0,0,0],
 [[0,1,2,3,4,5,6],[1,1,2,3,4,5,6],[2,2,1,2,3,4,5],
 Edit Distance: 3
```

```
import nltk

def find_minimum_edit_distance(word1
                               word2):

    distance = nltk.edit_distance(word1,
                                  word2)
    return distance
```

```
# Example

usageword1=
"kitten"word2
="sitting"
min_edit_dist
ance=find_mi
nimum_edit_
distance(word
1,word2)

print(f"The minimum edit distance between'{word1}'and'{word2}'is:{min_edit_distance}")
```

**Output:**

The minimum edit distance between Kitten and Sitting is: 3

## EXPERIMENT-4 : Implement n-gram language model

To create an N-gram language model using NLTK in Python, we'll use NLTK's ngrams function to generate n-grams from a text corpus. An N-gram language model predicts the probability of the next word in a sequence given the previous (n-1) words. Here's a Python program that demonstrates how to build an N-gram language model using NLTK.

```
import nltk
nltk.download('brown')

nltk.download('punkt')

from nltk.corpus import brown

from nltk.util import bigrams

from nltk.lm.preprocessing import pad_both_ends, padded_everygram_pipeline

from nltk.lm import MLE
```

```
corpus = brown.sents(categories="news")

test_sentence = ['There', "wasn't", 'a', 'bit', 'of', 'trouble', 'in', 'Texas']

test_sentence_bigrams = list(bigrams(pad_both_ends(test_sentence, n=2)))

print(test_sentence_bigrams)

train_data, vocab = padded_everygram_pipeline(2, corpus)

lm = MLE(2).fit(train_data, vocab)

print("Number of words in vocabulary is:", len(lm.vocab))

print(lm.counts)

prob = 1

for t in test_sentence_bigrams:

    score = lm.score(t[1], [t[0]])

    print(score)

    prob *= score

print(prob)
```

**OutPut:-**

```
[nltk_data] Downloading package brown to
[nltk_data]   /Users/varshuraodugyala/nltk_data...
[nltk_data] Package brown is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data]   /Users/varshuraodugyala/nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

```
[('<s>', 'There'), ('There', "wasn't"), ("wasn't", 'a'), ('a', 'bit'), ('bit', 'of'), ('of', 'trouble'), ('trouble', 'in'),
('in', 'Texas'), ('Texas', '</s>')]
```

```
Number of words in vocabulary is: 14397
<NgramCounter with 2 ngram orders and 214977 ngrams>
0.011464417045208739
0.017241379310344827
0.3333333333333333
0.002508780732563974
0.2857142857142857
0.001053001053001053
0.375
0.001584786053882726
0.125
3.6943506664397765e-15

0.125
```

## **EXPERIMENT-5 : Implement Naïve Bayes classification for sentiment analysis.**

To implement a Naive Bayes classifier for sentiment analysis using NLTK (Natural Language Toolkit) in Python, we'll use a dataset of movie reviews as an example. We'll preprocess the data, extract features, train the classifier, and then evaluate its performance.

```
import nltk
nltk.download('movie_reviews')

from nltk.corpus import movie_reviews
import random

documents = [(list(movie_reviews.words(fileid)), category)
             for category in movie_reviews.categories()
             for fileid in movie_reviews.fileids(category)] 

random.shuffle(documents)

all_words = nltk.FreqDist(w.lower() for w in movie_reviews.words())

word_features = list(all_words)[:2000]

def document_features(document):

    document_words = set(document)

    features = {}

    for word in word_features:

        features['contains({})'.format(word)] = (word in document_words)
```

```
return features

featuresets = [(document_features(d), c) for (d, c) in documents]

train_set, test_set = featuresets[100:], featuresets[:100]

classifier = nltk.NaiveBayesClassifier.train(train_set)

confusion_matrix = {"tp": 0, "tn": 0, "fp": 0, "fn": 0}

for i in range(len(test_set)):

    predicted = classifier.classify(test_set[i][0])

    actual = test_set[i][1] print(predicted, actual)

    if predicted == 'pos' and actual == 'pos':

        confusion_matrix['tp'] += 1

    elif predicted == 'neg' and actual == 'neg':

        confusion_matrix['tn'] += 1

    elif predicted == 'pos' and actual == 'neg':

        confusion_matrix['fp'] += 1

    else:

        confusion_matrix['fn'] += 1

print("\nConfusion Matrix:")

print("True Positives (TP): ", confusion_matrix["tp"])

print("True Negatives (TN): ", confusion_matrix["tn"])

print("False Positives (FP): ", confusion_matrix["fp"])

print("False Negatives (FN): ", confusion_matrix["fn"])
```

#### Output:

```
[nltk_data] Downloading package movie_reviews to
[nltk_data]   /Users/varshuraodugyala/nltk_data...
[nltk_data] Package movie_reviews is already up-to-date!
```

```
pos pos
neg neg
pos pos
neg neg
neg pos
```

neg neg  
pos pos  
neg neg  
neg neg  
neg neg  
pos pos  
pos pos  
pos pos  
neg neg  
pos pos  
neg neg  
neg neg  
neg neg  
pos pos  
neg pos  
pos pos  
neg neg  
pos pos  
pos neg  
neg pos  
neg pos  
neg neg  
neg neg  
neg pos  
pos pos  
pos neg  
neg neg  
neg pos  
pos pos  
pos neg  
pos pos  
neg neg  
neg neg  
pos pos  
pos pos  
neg neg  
neg neg  
neg neg  
pos pos  
pos pos  
neg neg  
neg neg  
pos neg

---

```
neg neg
pos neg
neg neg
neg neg
pos pos
pos neg
neg neg
pos pos
neg neg
pos pos
pos pos
neg neg
neg neg
neg neg
neg neg
pos pos
neg neg
```

Confusion Matrix:

True Positives (TP): 33  
True Negatives (TN): 44  
False Positives (FP): 13  
False Negatives (FN): 10

## EXPERIMENT-6 : Implement POS tagging using HMM

To implement a Hidden Markov Model (HMM) Part-of-Speech (POS) tagger for the Brown corpus using NLTK (version 3.x), we'll utilize NLTK's Hidden Markov Model Tagger class. However, as of NLTK version 3.x, direct support for HMM-based POS tagging using Hidden Markov Model Tagger is not available. Instead, NLTK provides an interface for using third-party implementations like the `hmm_tagger` from the `nltk.tag.hmm` module.

### 6.a)

```
import nltk
from nltk.corpus import brown
from nltk.tag import hmm
nltk.download('punkt')
```

```
sentences = brown.tagged_sents()  
  
trainer=hmm.HiddenMarkovModelTrainer()  
  
tagger = trainer.train(sentences)  
  
text= " this is a sample sentence for POS tagging in python"  
  
words=nltk.word_tokenize(text)  
  
tags=tagger.tag(words)  
  
for word,tag in tags:print(f'{word}: {tag}')
```

**Output :**

this:DT

is:BEZ

a:AT

sample:NN

sentence:NN

for:IN

POS:AT

tagging:AT

in:AT

python:AT

**6.b)**

```
import nltk  
  
from nltk.corpus import brown
```

```
from nltk.tag import hmm  
  
nltk.download('brown')  
  
brown_tagged_sentences = brown.tagged_sents(categories='news')  
  
size=int(len(brown_tagged_sentences)*0.9)  
  
train_sentences = brown_tagged_sentences[:size]  
  
test_sentences = brown_tagged_sentences[:size]  
  
trainer = hmm.HiddenMarkovModelTrainer()  
  
tagger=trainer.train(train_sentences)  
  
print(tagger.accuracy(test_sentences))
```

Output :

```
[nltk_data] Downloading package brown to /root/nltk_data...
```

```
[nltk_data] Package brown is already up-to-date!
```

```
0.98089945979386
```

## **EXPERIMENT-7 : Implementation of Implement CKY parsing algorithm**

The CKY (Cocke-Kasami-Younger) algorithm is a parsing algorithm for context-free grammars typically used in natural language processing. Below is an implementation of the CKY parser in Python

```
def print_chart(chart, n):  
  
    for i in range(n):  
  
        for j in range(n+1):  
  
            print(chart[i][j], end="\t")  
  
        print()
```

```
def CKY_PARSE(words, grammar):
    n = len(words)

    table = [[set() for _ in range(n+1)] for _ in range(n+1)]

    # Fill diagonal with preterminal rules
    for i in range(n):
        word = words[i]
        for lhs, rhs in grammar:
            if rhs == (word,):
                table[i][i+1].add(lhs)

    # Fill upper cells using binary rules
    for length in range(2, n+1):
        for i in range(n - length + 1):
            j = i + length
            for k in range(i+1, j):
                for lhs, rhs in grammar:
                    if len(rhs) == 2:
                        if rhs[0] in table[i][k] and rhs[1] in table[k][j]:
                            table[i][j].add(lhs)
    return table

sentence = "the dog chased the cat" words = sentence.split()

n = len(words)

grammar = [ ('S', ('NP', 'VP')), ('NP', ('DET', 'NOMINAL')), ('VP', ('VERB', 'NP')), ('NOMINAL', ('cat',)), ('NOMINAL', ('dog',)), ('VERB', ('chased',)), ('DET', ('the',)) ]

chart = CKY_PARSE(words, grammar)

print(words, "\n")

print_chart(chart, n)

start_symbol = 'S'

if start_symbol in chart[0][n]:
    print("The sentence is grammatically correct.")

else:
    print("The sentence is not grammatically correct.")
```

**Output:-**

```
['the', 'dog', 'chased', 'the', 'cat']

set() {'DET'} {'NP'} set() set() {'S'}

set() set() {'NOMINAL'} set() set() set()

set() set() set() {'VERB'} set() set()

set() set() set() set() {'DET'} {'NP'}

set() set() set() set() set() {'NOMINAL'}

set() set() set() set() set()

The sentence is grammatically correct.
```

**EXPERIMENT-8 : Implement PCKY parsing algorithm**

The Probabilistic CKY (PCKY) parser is an extension of the CKY parser that incorporates probabilities for parsing probabilistic context-free grammars (PCFGs). Below is an implementation of a PCKY parser in Python

```
def print_chart(chart, n, non_terminals):

    for p in range(n+1):

        for q in range(n+1):

            print('[', p, ',', q, ']:', end=" ")

            for nt in non_terminals:

                if chart[p][q][nt] > 0:

                    print('{', nt, ':', chart[p][q][nt], '}', end="")
```

```
    print()

    print()

def PCKY_PARSE(words, grammar, non_terminals):

    n = len(words)

    print(words, "\n")

    table = [[dict() for _ in range(n+1)] for _ in range(n+1)]
    for i in range(n+1):
        for j in range(n+1):
            for nt in non_terminals:
                table[i][j][nt] = 0.0

    for j in range(1, n+1):
        for lhs, rhs, pr in grammar:
            if rhs == (words[j-1],):
                table[j-1][j][lhs] = pr

    # Fill upper cells
    for i in range(j-2, -1, -1):
        for k in range(i+1, j):
            for lhs, rhs, pr in grammar:
                if len(rhs) == 2 and table[i][k][rhs[0]] > 0
and table[k][j][rhs[1]] > 0:
                    prob = pr *
table[i][k][rhs[0]]*table[k][j][rhs[1]]
                    if table[i][j][lhs] < prob:
                        table[i][j][lhs] = prob

    return table

sentence = "the flight includes a meal"

words = sentence.split()

n = len(words)

grammar = [ ('S', ('NP', 'VP'), 0.80), ('NP', ('DET', 'NOMINAL'), 0.30),
('VP', ('VERB', 'NP'), 0.20), ('NOMINAL', ('meal',), 0.01), ('NOMINAL',
```

```
('flight',), 0.02), ('VERB', ('includes',), 0.05), ('DET', ('the',),  
0.40), ('DET', ('a',), 0.40) ]  
  
non_terminals = ['S', 'NP', 'VP', 'DET', 'NOMINAL', 'VERB']  
  
chart = PCKY_PARSE(words, grammar, non_terminals)  
  
print_chart(chart, n, non_terminals)  
  
start_symbol = 'S'  
  
if chart[0][n]['S'] > 0:  
    print("The sentence is grammatically correct.")  
  
else:  
  
    print("The sentence is not grammatically correct.")
```

**OUTPUT:**

['the', 'flight', 'includes', 'a', 'meal'] [ 0 , 0 ]:

[ 0 , 1 ] : { DET : 0.4 }

[ 0 , 2 ] : { NP : 0.0024 } [ 0 , 3 ] :

[ 0 , 4 ] :

[ 0 , 5 ] : { S : 2.3040000000000003e-08 }

[ 1 , 0 ] :

[ 1 , 1 ] :

[ 1 , 2 ] : { NOMINAL : 0.02 } [ 1 , 3 ] :

[ 1 , 4 ] :

[ 1 , 5 ] :

[ 2 , 0 ]:  
[ 2 , 1 ]:  
[ 2 , 2 ]:  
[ 2 , 3 ]: { VERB : 0.05 } [ 2 , 4 ]:  
[ 2 , 5 ]: { VP : 1.2000000000000002e-05 } [ 3 , 0 ]:  
[ 3 , 1 ]:  
[ 3 , 2 ]:  
[ 3 , 3 ]:  
[ 3 , 4 ]: { DET : 0.4 }  
[ 3 , 5 ]: { NP : 0.0012 } [ 4 , 0 ]:  
[ 4 , 1 ]:  
[ 4 , 2 ]:  
[ 4 , 3 ]:  
[ 4 , 4 ]:  
[ 4 , 5 ]: { NOMINAL : 0.01 } [ 5 , 0 ]:  
[ 5 , 1 ]:  
[ 5 , 2 ]:  
[ 5 , 3 ]:  
[ 5 , 4 ]:  
[ 5 , 5 ]:

The sentence is grammatically correct.

## **EXPERIMENT-9: Implementation of Computing cosine similarity between the words using term document matrix and term-term matrix**

Cosine similarity is a metric used to measure how similar two vectors are, regardless of their size. It is often used in natural language processing and information retrieval to assess the similarity between documents.

```
import nltk
import random
import math
from nltk.corpus import brown, stopwords
nltk.download('brown')
nltk.download('stopwords')

def extract_words(document):
    all_terms_list = brown.words(fileids=document)
    only_words_list = [w.lower() for w in all_terms_list if w.isalpha()]
    stopwords_list = stopwords.words('english')
    final_terms_list = [w for w in only_words_list if w not in stopwords_list]
    return final_terms_list

def freq(word, document):
    d_terms = extract_words(document)
```

```
fdist = nltk.FreqDist(d_terms)

return fdist[word]

doc_names = ['ca01', 'ca02', 'ca03', 'ca04']

vocab = set()

for doc in doc_names:

    vocab.update(extract_words(doc))

vocab_len = len(vocab)

print("Length of vocabulary =", vocab_len)

word1 = list(vocab)[random.randint(0, len(vocab)-1)]

word2 = list(vocab)[random.randint(0, len(vocab)-1)]

print("word-1:", word1)

print("word-2:", word2)

word1_vector = [freq(word1, doc) for doc in doc_names]

word2_vector = [freq(word2, doc) for doc in doc_names]

print("word-1-vector:", word1_vector)

print("word-2-vector:", word2_vector)

dot_product = sum(word1_vector[i] * word2_vector[i] for i in range(len(doc_names)))

vector1_len = math.sqrt(sum(w2 for w in word1_vector))

vector2_len = math.sqrt(sum(w2 for w in word2_vector))

cos_theta = dot_product / (vector1_len * vector2_len) if vector1_len*vector2_len != 0 else 0
print(f"cos_theta({word1}, {word2}) =", cos_theta)
```

**OUTPUT:**

length of vocabulary = 2023

word-1: ignored

```
word-2: interest  
word-1-vector: [0, 0, 1, 0]  
word-2-vector: [2, 0, 0, 0]  
cos_theta( ignored , interest ) = 0.0
```

## **EXPERIMENT-10: Implementation of TF-IDF matrix for the given document set**

To calculate cosine similarity using TF-IDF (Term Frequency-Inverse Document Frequency), you can use the Scikit-learn library, which provides tools for computing TF-IDF vectors and then calculating cosine similarity between them.

```
import nltk  
import random  
import math  
from nltk.corpus import brown, stopwords  
nltk.download('brown')  
nltk.download('stopwords')  
doc_names = ['ca01', 'ca02', 'ca03', 'ca04']  
  
def extract_words(document):  
    all_terms_list = brown.words(fileids=document)  
    only_words_list = [w.lower() for w in all_terms_list if w.isalpha()]  
    stopwords_list = stopwords.words('english')
```

```
final_terms_list = [w for w in only_words_list if w not in stopwords_list]

return final_terms_list

def term_freq(word, document):

    d_terms=extract_words(document)

    fdist = nltk.FreqDist(d_terms)

    return math.log10(fdist[word] + 1)

def idf(word):

    df=0 for doc in doc_names:

        if word in extract_words(doc):

            df += 1

    return math.log10(len(doc_names) / df) if df != 0 else 0

vocab = set()

for doc in doc_names:

    vocab.update(extract_words(doc))

vocab_len = len(vocab)

print("Length of vocabulary =", vocab_len)

word1 = list(vocab)[random.randint(0, len(vocab)-1)]

word2 = list(vocab)[random.randint(0, len(vocab)-1)]

print("word-1:", word1)

print("word-2:", word2)

word1_vector = [term_freq(word1, doc) * idf(word1) for doc in doc_names]

word2_vector = [term_freq(word2, doc) * idf(word2) for doc in doc_names]

print("word-1-vector:", word1_vector)

print("word-2-vector:", word2_vector)

dot_product = sum(word1_vector[i] * word2_vector[i] for i in range(len(doc_names)))

vector1_len = math.sqrt(sum(w2 for w in word1_vector))

vector2_len = math.sqrt(sum(w2 for w in word2_vector))
```

```
cos_theta = dot_product / (vector1_len * vector2_len) if vector1_len * vector2_len != 0 else 0
print(f"cos_theta({{word1}}, {{word2}}) =", cos_theta)
```

**OUTPUT:**

Length	of	vocabulary	=	2023
word-1:				opelika
word-2:				compulsory
word-1-vector:	[0.1812381165789131,	0.0,	0.0,	0.0]
word-2-vector:	[0.0,	0.0,	0.1812381165789131,	0.0]
cos_theta(opelika, compulsory) = 0.0				

**EXPERIMENT-11: Language Model Using Feed forward Neural Network**

Language modeling using a Feed-Forward Neural Network (FFNN) can be implemented with libraries like Tensor Flow or PyTorch. In this example, I'll use Tensor Flow to create a simple language model. The model will be trained to predict the next word in a sequence of words.

```
import tensorflow as tf

from tensorflow.keras.preprocessing.text import Tokenizer

from tensorflow.keras.preprocessing.sequence import pad_sequences

import numpy as np

corpus = [ 'This is a simple example', 'Language modeling is interesting', 'Neural networks are
powerful', 'Feed-forward networks are common in natural language processing' ]

tokenizer = Tokenizer() tokenizer.fit_on_texts(corpus)

total_words = len(tokenizer.word_index) + 1

input_sequences = []

for line in corpus:

    token_list = tokenizer.texts_to_sequences([line])[0]

    for i in range(1, len(token_list)):

        n_gram_sequence = token_list[:i+1]
```

```
        input_sequences.append(n_gram_sequence)

max_sequence_length = max([len(x) for x in input_sequences])

input_sequences = pad_sequences(input_sequences, maxlen=max_sequence_length, padding='pre')

X, y = input_sequences[:, :-1], input_sequences[:, -1]

y = tf.keras.utils.to_categorical(y, num_classes=total_words)

model = tf.keras.Sequential([
    tf.keras.layers.Embedding(total_words, 50,
    input_length=max_sequence_length-1),
    tf.keras.layers.LSTM(100),
    tf.keras.layers.Dense(total_words, activation='softmax') ])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model.fit(X, y, epochs=100, verbose=1)

seed_text = "Neural networks" next_words = 7 # Number of words to generate

for _ in range(next_words):

    token_list = tokenizer.texts_to_sequences([seed_text])[0]

    token_list = pad_sequences([token_list], maxlen=max_sequence_length-1, padding='pre')
    predicted_index = np.argmax(model.predict(token_list, verbose=0), axis=-1)[0]

    output_word = ""
    for word, index in tokenizer.word_index.items():
        if index == predicted_index:
            output_word = word
            break

    seed_text += " " + output_word

print(seed_text)
```

**OUTPUT:**

Epoch 1/1001/1 [=====] - ETA: 0s - loss: 2.8825 -  
accuracy:0.1250

---

```
1/1 [=====] - 1s 897ms/step - loss: 2.8825 - accuracy: 0.1250
Epoch 100/100 1/1 [=====] - ETA: 0s - loss: 0.0587 -
accuracy:1.0000
1/1 [=====] - 0s 8ms/step - loss: 0.0587 - accuracy: 1.0000
Neural networks are powerful in Natural Language Processing.
```

## EXPERIMENT-12: Implement Language Model Using RNN

Language modeling using a Recurrent Neural Network (RNN) can be implemented effectively using libraries like TensorFlow. Here's a step-by-step guide to building a language model using an RNN

```
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
import numpy as np
# Toy dataset corpus = ['This is a simple example',
                      'Language modeling is interesting',
                      'Neural networks are powerful',
                      'Recurrent neural networks capture sequences well'] # Tokenize the text
tokenizer = Tokenizer()
tokenizer.fit_on_texts(corpus)
total_words = len(tokenizer.word_index) + 1
# Create input sequences and labels
input_sequences = []
for line in corpus:
```

```
token_list = tokenizer.texts_to_sequences([line])[0]

for i in range(1, len(token_list)):

    n_gram_sequence = token_list[:i+1]

    input_sequences.append(n_gram_sequence)

max_sequence_length = max([len(x) for x in input_sequences])

input_sequences = pad_sequences(input_sequences, maxlen=max_sequence_length, padding='pre')

X, y = input_sequences[:, :-1], input_sequences[:, -1]

y = tf.keras.utils.to_categorical(y, num_classes=total_words) # Build the model

model = tf.keras.Sequential([
    tf.keras.layers.Embedding(total_words, 50, input_length=max_sequence_length-1),
    tf.keras.layers.LSTM(100),
    tf.keras.layers.Dense(total_words, activation='softmax')])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(X, y, epochs=100, verbose=1)

# Generate text using the trained model

seed_text = "Recurrent neural networks" next_words = 5

for _ in range(next_words):

    token_list = tokenizer.texts_to_sequences([seed_text])[0]

    token_list = pad_sequences([token_list], maxlen=max_sequence_length-1, padding='pre')

    predicted = np.argmax(model.predict(token_list), axis=-1)

    output_word = ""

    for word, index in tokenizer.word_index.items():

        if index == predicted:

            output_word = word

            break

    seed_text += " " + output_word

print(seed_text)
```

**OUTPUT:**

Epoch 1/100 1/1 [=====] - ETA: 0s - loss: 2.8342 -  
accuracy:0.0667

1/1 [=====] - 1s 1s/step - loss: 2.8342 - accuracy: 0.0667  
Epoch 100/100 1/1 [=====] - ETA: 0s - loss: 0.3300 -  
accuracy:1.0000

1/1 [=====] - 0s 0s/step - loss: 0.3300 - accuracy: 1.0000

1/1 [=====] - 0s 0s/step - loss: 0.3300 - accuracy: 1.0000  
1/1 [=====] - ETA:0s

1/1 [=====] - 0s 336ms/step

1/1 [=====] - ETA:  
0s

1/1 [=====] - 0s 16ms/step

1/1 [=====] - ETA:  
0s

1/1 [=====] - 0s 8ms/step

1/1 [=====] - ETA:  
0s

1/1 [=====] - 0s 16ms/step

Recurrent neural networks capture sequences well well well

## **EXPERIMENT-13: Implement Perform Text Analytics**

Text classification is a process in natural language processing (NLP) where text is categorized into predefined classes or labels.

Sentiment Analysis: Determining whether a piece of text expresses positive, negative, or neutral sentiment.

```
from sklearn.feature_extraction.text import CountVectorizer  
from sklearn.naive_bayes import MultinomialNB  
  
# Sample data  
documents = ["I love programming.", "Python is great for NLP.", "Text analytics is fun!"]  
labels = [1, 1, 0] # 1 for positive, 0 for neutral/negative  
  
# Vectorization  
vectorizer = CountVectorizer()  
X = vectorizer.fit_transform(documents)
```

---

```
# Model training  
clf = MultinomialNB()  
clf.fit(X, labels)  
  
# Prediction  
test_docs = ["I enjoy coding.", "NLP is interesting."  
X_test = vectorizer.transform(test_docs)  
predictions = clf.predict(X_test)  
print(predictions)
```

**OutPut:**

[1 1 ]

## ADDITIONAL PROGRAMS

## 1) AMBIGUITY

Python program:

```
pip install nltk
```

```
from nltk import CFG, ChartParser
```

```
# -----
```

```
# Example 1: Lexical Ambiguity
```

```
# -----
```

```
word = "bank"
```

```
meanings = [
```

```
    "A financial institution where people deposit and withdraw money",
```

```
    "The land alongside a river",
```

```
    "To tilt an airplane while turning"
```

```
]
```

```
print("LEXICAL AMBIGUITY:")
```

```
print(f"The word '{word}' has {len(meanings)} meanings:")
```

```
for i, meaning in enumerate(meanings, 1):
```

```
    print(f'{i}. {meaning}')
```

```
print("\n" + "-"*50 + "\n")
```

```
# -----
```

```
# Example 2: Syntactic Ambiguity
```

```
# -----
# Grammar to parse "I saw the man with a telescope"
grammar = CFG.fromstring(""""
S -> NP VP
VP -> V NP | VP PP
PP -> P NP
NP -> Det N | NP PP | 'T'
V -> 'saw'
Det -> 'the' | 'a'
N -> 'man' | 'telescope'
P -> 'with'
""")
```

```
sentence = ['T', 'saw', 'the', 'man', 'with', 'a', 'telescope']
parser = ChartParser(grammar)
```

```
print("SYNTACTIC AMBIGUITY:")
print("Sentence: 'I saw the man with a telescope'")
print("Possible parses:\n")
```

```
for tree in parser.parse(sentence):
    print(tree)      # Print the parse tree in text form
    tree.pretty_print() # Pretty print the tree
    print()
```

## Output:

## LEXICAL AMBIGUITY:



The word 'bank' has 3 meanings:

1. A financial institution where people deposit and withdraw money
  2. The land alongside a river
  3. To tilt an airplane while turning
- 

## SYNTACTIC AMBIGUITY:

Sentence: 'I saw the man with a telescope'

Possible parses:



I saw the man with a telescope

## 2)Text classification program:

```
pip install scikit-learn  
import spacy from spacy import displacy  
nlp = spacy.load("en_core_web_sm")  
  
text = """ Apple Inc. is looking at buying U.K. startup for $1 billion. Tim Cook, the CEO of Apple,  
said this deal will be completed by next month. Barack Obama was the 44th president of the United  
States. """  
  
doc = nlp(text)  
  
for ent in doc.ents: print(f"Text: {ent.text} | Label: {ent.label_}")  
  
displacy.render(doc, style="ent", jupyter=True)
```

## OUTPUT:

```
Text: Apple Inc. | Label: ORG  
Text: U.K. | Label: GPE  
Text: $1 billion | Label: MONEY  
Text: Tim Cook | Label: PERSON  
Text: Apple | Label: ORG  
Text: next month | Label: DATE  
Text: Barack Obama | Label: PERSON  
Text: 44th | Label: ORDINAL  
Text: the United States | Label: GPE
```

Apple Inc. **ORG** is looking at buying U.K. **GPE** startup for \$1 billion **MONEY** .  
Tim Cook **PERSON** , the CEO of Apple **ORG** , said this deal will be completed by next month  
**DATE** .  
Barack Obama **PERSON** was the 44th **ORDINAL** president of the United States **GPE** .

### 3)Smoothing

```
pip install sentence-transformers scikit-learn

from sentence_transformers import SentenceTransformer from sklearn.metrics.pairwise import
cosine_similarity import numpy as np

model = SentenceTransformer('all-MiniLM-L6-v2') # small, fast, good accuracy

sentences = [ "I love playing football.", "Soccer is my favorite sport.", "The weather is sunny
today.", "It is raining outside." ]

sentence_vectors = model.encode(sentences)

print("Sentence embedding shape:", sentence_vectors.shape)

similarity_matrix = cosine_similarity(sentence_vectors)

print("\nCosine Similarity Matrix:\n") print(similarity_matrix)

most_similar_idx = np.argmax(similarity_matrix[0][1:]) + 1 # ignore self-similarity

print(f"\nSentence most similar to '{sentences[0]}' is '{sentences[most_similar_idx]}'"")
```

#### OUTPUT:

Sentence embedding shape: (4, 384)

Cosine Similarity Matrix:

```
[[ 1.      0.65746427 -0.02304317  0.01971502]
 [ 0.65746427  1.0000002   0.05700713  0.04840674]
 [-0.02304317  0.05700713  1.0000001   0.46222505]
 [ 0.01971502  0.04840674  0.46222505  0.99999976]]
```

Sentence most similar to 'I love playing football.' is 'Soccer is my favorite sport.'

### 4)Text summarization:

```
pip install transformers torch

from transformers import pipeline
```

```
summarizer = pipeline("summarization", model="t5-small") # or "facebook/bart-large-cnn"

text = """ Natural Language Processing (NLP) is a field of artificial intelligence that focuses on the interaction between computers and humans using natural language. The ultimate objective of NLP is to read, decipher, understand, and make sense of human languages in a manner that is valuable. Many challenges in NLP involve natural language understanding, natural language generation, sentiment analysis, text summarization, machine translation, and question answering. """

summary = summarizer(text, max_length=50, min_length=25, do_sample=False)

print("Original Text:\n", text)
print("\nSummarized Text:\n", summary[0]['summary_text'])

print("\nSummarized Text:\n", summary[0]['summary_text'])
```

## OUTPUT:

### Original Text:

Natural Language Processing (NLP) is a field of artificial intelligence that focuses on the interaction between computers and humans using natural language.  
The ultimate objective of NLP is to read, decipher, understand, and make sense of human languages in a manner that is valuable.  
Many challenges in NLP involve natural language understanding, natural language generation, sentiment analysis, text summarization, machine translation, and question answering.

### Summarized Text:

natural language processing (NLP) is a field of artificial intelligence that focuses on the interaction between computers and humans using natural language . many challenges in NLP involve natural language understanding, natural language generation, sentiment analysis, text summarization, machine translation .

## 5)RECOMMENDATION SYSTEM:

```
from sklearn.feature_extraction.text
import TfidfVectorizer from sklearn.metrics.pairwise
import cosine_similarity
```

```
movies = [ {"title": "The Matrix", "description": "A hacker learns about the true nature of reality and fights against controllers."}, {"title": "Inception", "description": "A thief enters dreams to plant an idea into a CEO's mind."}, {"title": "Interstellar", "description": "Explorers travel through a wormhole in space to ensure humanity's survival."}, {"title": "The Lord of the Rings", "description": "A hobbit journeys to destroy a powerful ring and save Middle-earth."}, {"title": "The Social Network", "description": "The story of the founding of Facebook and the legal battles that followed."} ]  
  
descriptions = [movie['description'] for movie in movies]  
  
vectorizer = TfidfVectorizer(stop_words='english') tfidf_matrix = vectorizer.fit_transform(descriptions)  
  
similarity_matrix = cosine_similarity(tfidf_matrix)  
  
def recommend(title, movies=movies, similarity_matrix=similarity_matrix, top_n=3): # Find index of the movie  
  
    idx = next(i for i, m in enumerate(movies) if m['title'] == title)  
  
    # Get similarity scores for this movie  
    sim_scores = list(enumerate(similarity_matrix[idx]))  
  
    # Sort by similarity and get top_n (excluding the movie itself)  
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)  
    sim_scores = sim_scores[1:top_n+1]  
  
    # Print recommendations  
    print(f"Movies similar to '{title}':")  
    for i, score in sim_scores:  
        print(f"- {movies[i]['title']} (Similarity: {score:.2f})")  
recommend("Inception")
```

### Output:

```
Movies similar to 'Inception':  
- The Matrix (Similarity: 0.00)  
- Interstellar (Similarity: 0.00)  
- The Lord of the Rings (Similarity: 0.00)
```

### 6) Encoder decoder machine translation:

```
import numpy as np
```

```
from tensorflow.keras.models import Model

from tensorflow.keras.layers import Input, LSTM, Embedding, Dense

from tensorflow.keras.preprocessing.text import Tokenizer

from tensorflow.keras.preprocessing.sequence import pad_sequences

input_texts = ["hello", "how are you", "i am fine", "thank you", "good morning"]

target_texts = [" bonjour ", " comment ça va ", " je vais bien ", " merci ", " bonjour "]

input_tokenizer = Tokenizer()

input_tokenizer.fit_on_texts(input_texts)

input_sequences = input_tokenizer.texts_to_sequences(input_texts)

max_encoder_seq_length = max(len(seq) for seq in input_sequences)

num_encoder_tokens = len(input_tokenizer.word_index) + 1

encoder_input_data = pad_sequences(input_sequences, maxlen=max_encoder_seq_length,
padding='post')

target_tokenizer = Tokenizer(filters="")

target_tokenizer.fit_on_texts(target_texts)

target_sequences = target_tokenizer.texts_to_sequences(target_texts)

max_decoder_seq_length = max(len(seq) for seq in target_sequences)

num_decoder_tokens = len(target_tokenizer.word_index) + 1

decoder_input_data = pad_sequences([seq[:-1] for seq in target_sequences],
maxlen=max_decoder_seq_length-1, padding='post')

decoder_target_data = pad_sequences([seq[1:] for seq in target_sequences],
maxlen=max_decoder_seq_length-1, padding='post')

decoder_target_data_oh = np.zeros((len(input_texts), max_decoder_seq_length-1,
num_decoder_tokens), dtype='float32')

for i, seq in enumerate(decoder_target_data):

    for t, word_idx in enumerate(seq):

        if word_idx > 0:
```

```

decoder_target_data_oh[i, t, word_idx] = 1.0

latent_dim = 32

embedding_dim = 32

encoder_inputs = Input(shape=(max_encoder_seq_length,))

enc_emb = Embedding(num_encoder_tokens, embedding_dim,
input_length=max_encoder_seq_length)(encoder_inputs)

encoder_lstm = LSTM(latent_dim, return_state=True)

encoder_outputs, state_h, state_c = encoder_lstm(enc_emb) encoder_states = [state_h, state_c]

decoder_inputs = Input(shape=(max_decoder_seq_length-1,))

dec_emb = Embedding(num_decoder_tokens, embedding_dim,
input_length=max_decoder_seq_length-1)(decoder_inputs)

decoder_lstm = LSTM(latent_dim, return_sequences=True, return_state=True)

decoder_outputs, _, _ = decoder_lstm(dec_emb, initial_state=encoder_states)

decoder_dense = Dense(num_decoder_tokens, activation='softmax')

decoder_outputs = decoder_dense(decoder_outputs)

model = Model([encoder_inputs, decoder_inputs], decoder_outputs)

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model.summary()


model.fit([encoder_input_data, decoder_input_data], decoder_target_data_oh, batch_size=2,
epochs=50, verbose=1)

```

## OUTPUT:

**Model: "functional"**

Layer (type)	Output Shape	Param #	Connected to
input_layer_1	(None, 3)	0	-



**Total params:** 17,707 (69.17 KB)

**Trainable params:** 17,707 (69.17 KB)

**Non-trainable params:** 0 (0.00 B)

Epoch 1/50

3/3 2s 7ms/step - accuracy: 0.1500 - loss: 1.6760

Epoch 2/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.6684

Epoch 3/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.6608

Epoch 4/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.6526

Epoch 5/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.6446

Epoch 6/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.6355

Epoch 7/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.6247

Epoch 8/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.6127  
Epoch 9/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.5992  
Epoch 10/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.5839  
Epoch 11/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.5649  
Epoch 12/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.5423  
Epoch 13/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.5158  
Epoch 14/50

3/3 0s 7ms/step - accuracy: 0.2500 - loss: 1.4838  
Epoch 15/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.4497  
Epoch 16/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.4098  
Epoch 17/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.3674  
Epoch 18/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.3398  
Epoch 19/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.3221  
Epoch 20/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.3089  
Epoch 21/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.2993  
Epoch 22/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.2899  
Epoch 23/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.2804  
Epoch 24/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.2668  
Epoch 25/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.2504  
Epoch 26/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.2313  
Epoch 27/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.2096  
Epoch 28/50

3/3 0s 5ms/step - accuracy: 0.2500 - loss: 1.1874  
Epoch 29/50

3/3 0s 6ms/step - accuracy: 0.2500 - loss: 1.1644  
Epoch 30/50

```
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.2500 - loss: 1.1467
Epoch 31/50
3/3 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.2500 - loss: 1.1232
Epoch 32/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.2500 - loss: 1.1045
Epoch 33/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.2500 - loss: 1.0826
Epoch 34/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.2500 - loss: 1.0600
Epoch 35/50
3/3 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.2500 - loss: 1.0368
Epoch 36/50
3/3 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.2500 - loss: 1.0128
Epoch 37/50
3/3 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.3000 - loss: 0.9888
Epoch 38/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.3500 - loss: 0.9661
Epoch 39/50
3/3 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.3500 - loss: 0.9444
Epoch 40/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.3500 - loss: 0.9232
Epoch 41/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.3500 - loss: 0.9039
Epoch 42/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.3500 - loss: 0.8851
Epoch 43/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.3500 - loss: 0.8670
Epoch 44/50
3/3 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.4000 - loss: 0.8505
Epoch 45/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.4000 - loss: 0.8341
Epoch 46/50
3/3 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.4000 - loss: 0.8168
Epoch 47/50
3/3 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.4500 - loss: 0.8035
Epoch 48/50
3/3 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.5000 - loss: 0.7849
Epoch 49/50
3/3 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.5000 - loss: 0.7692
Epoch 50/50
3/3 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.5000 - loss: 0.7505
```

<keras.src.callbacks.history.History at 0x14ea50860>

## 7)Named entity recognition:

```
pip install spacy  
!python -m spacy download en_core_web_sm  
  
import spacy nlp = spacy.load("en_core_web_sm")  
  
import spacy  
  
nlp = spacy.load("en_core_web_sm")  
  
text = """Apple is planning to open a new office in Hyderabad, India by 2026. Tim Cook, the CEO of  
Apple, met with Narendra Modi to discuss investment opportunities."""  
  
doc = nlp(text)  
  
print("Named Entities, their labels, and positions:\n")  
  
for ent in doc.ents:  
  
    print(f"{ent.text}<25} | {ent.label_}<10} | Start: {ent.start_char}, End: {ent.end_char}")  
  
print("\nEntity Label Explanation:") print(spacy.explain("GPE")) print(spacy.explain("ORG"))  
print(spacy.explain("PERSON"))
```

### Output:

Named Entities, their labels, and positions:

Apple	ORG	Start: 0, End: 5
Hyderabad	GPE	Start: 42, End: 51
India	GPE	Start: 53, End: 58
2026	DATE	Start: 62, End: 66
Tim Cook	PERSON	Start: 80, End: 88
Apple	ORG	Start: 101, End: 106
Narendra Modi	PERSON	Start: 117, End: 130

Entity Label Explanation:

Countries, cities, states

Companies, agencies, institutions, etc.

People, including fictional

## 8)Word to vec:

```
pip install gensim
pip install nltk
import nltk
nltk.download('punkt')
import nltk
nltk.download("punkt_tab")
from gensim.models import Word2Vec import nltk
nltk.download('punkt')

text = """Natural Language Processing is a field of Artificial Intelligence. It deals with how computers understand human language. Word embeddings like Word2Vec capture semantic meaning of words."""

sentences = nltk.sent_tokenize(text) data = [nltk.word_tokenize(sentence.lower()) for sentence in sentences]

model = Word2Vec(sentences=data, vector_size=50, window=5, min_count=1, sg=1)

print("Vector representation for 'language':\n", model.wv['language'])

print("\nWords most similar to 'language':") print(model.wv.most_similar('language'))

print("\nSimilarity between 'language' and 'computers':", model.wv.similarity('language', 'computers'))
```

## Output:

Vector representation for 'language':  
[-0.01631313 0.00898634 -0.00828216 0.00164936 0.01698417 -0.00893383  
 0.00903378 -0.01356238 -0.00709291 0.01878958 -0.00315817 0.00064451  
 -0.00826876 -0.01536609 -0.00303185 0.0049505 -0.00176973 0.01108032  
 -0.00550026 0.00451065 0.01092105 0.01671072 -0.00290536 -0.01840301  
 0.00874378 0.00115624 0.01488257 -0.00161443 -0.00528963 -0.01750265  
 -0.00171541 0.00567086 0.01080243 0.01410966 -0.01140788 0.00371402  
 0.01217452 -0.00960064 -0.00619748 0.01360284 0.00326445 0.00038038  
 0.00693825 0.00044166 0.01925861 0.01011735 -0.01783446 -0.01409558  
 0.00180856 0.01277417]

Words most similar to 'language':

[('understand', 0.23005631566047668), ('embeddings', 0.16102485358715057), ('word', 0.14890338480472565), ('deals', 0.12482792884111404), ('with', 0.08072882890701294), ('is', 0.07400191575288773), ('word2vec', 0.055481769144535065), ('!', 0.04347098991274834),

---

```
('processing', 0.018909117206931114), ('how', 0.01197921484708786)]
```

Similarity between 'language' and 'computers': -0.034299165

## 9)word disambiguity:

```
import nltk
from nltk.corpus import wordnet as wn
from nltk.wsd import lesk

# Download resources (only the first time)
nltk.download('punkt')
nltk.download('wordnet')
nltk.download('omw-1.4')

# Example sentences with the ambiguous word "bank"
sentences = [
    "He deposited money in the bank.",
    "The fisherman sat on the bank of the river."
]

for sent in sentences:
    tokens = nltk.word_tokenize(sent)
    # Apply Lesk algorithm for disambiguation
    sense = lesk(tokens, "bank")
    print(f"Sentences: {sent}")
    print(f"Predicted Sense: {sense}")
    print(f"Definition: {sense.definition() if sense else 'No definition found'}\n")
```

## Output:

```
Sentence: He deposited money in the bank.
Predicted Sense: Synset('savings_bank.n.02')
Definition: a container (usually with a slot in the top) for keeping money at home
```

Sentence: The fisherman sat on the bank of the river.  
Predicted Sense: Synset('bank.v.07')  
Definition: cover with ashes so to control the rate of burning

### 10)sentence to vec:

```
from sklearn.feature_extraction.text import TfidfVectorizer  
  
from sklearn.metrics.pairwise import cosine_similarity  
  
sentences = [ "I love playing football.", "Soccer is my favorite sport.", "The Eiffel Tower is in  
Paris." ]  
  
vectorizer = TfidfVectorizer()  
  
sentence_vectors = vectorizer.fit_transform(sentences)  
  
print("Vector shape:", sentence_vectors.shape)  
  
similarity = cosine_similarity(sentence_vectors[0], sentence_vectors[1])  
  
print("Similarity between first two sentences:", similarity[0][0])
```

### Output:

```
Vector shape: (3, 13)  
Similarity between first two sentences: 0.0
```

### 11)TEXT GENERATION WITH LSTM:

```
import numpy as np  
import tensorflow as tf  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import LSTM, Dense, Embedding  
  
# Sample text (toy dataset)
```

```

text = "lstm is a type of recurrent neural network. it is used in natural language processing."
chars = sorted(list(set(text)))
char_to_idx = {c: i for i, c in enumerate(chars)}
idx_to_char = {i: c for i, c in enumerate(chars)}

# Convert text to integers
encoded = [char_to_idx[c] for c in text]

# Prepare sequences
seq_length = 10
x = []
y = []
for i in range(len(encoded) - seq_length):
    x.append(encoded[i:i+seq_length])
    y.append(encoded[i+seq_length])
x = np.array(x)
y = np.array(y)

# Build model
model = Sequential()
model.add(Embedding(len(chars), 50, input_length=seq_length))
model.add(LSTM(100))
model.add(Dense(len(chars), activation='softmax'))

model.compile(loss='sparse_categorical_crossentropy', optimizer='adam')

# Train
model.fit(x, y, epochs=50, verbose=1)

# Generate text
seed = "lstm is a "
seed_encoded = [char_to_idx[c] for c in seed]
for _ in range(50):
    x_pred = np.array(seed_encoded[-seq_length:]).reshape(1, seq_length)
    pred = np.argmax(model.predict(x_pred, verbose=0))
    seed_encoded.append(pred)

generated_text = ''.join(idx_to_char[i] for i in seed_encoded)
print("Generated text:")
print(generated_text)

```

**OUTPUT:**

```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/embedding.py:97: UserWarning: Argument
`input_length` is deprecated. Just remove it.
  warnings.warn(

```

Epoch 1/50  
3/3 ————— **4s** 50ms/step - loss: 3.0434  
Epoch 2/50  
3/3 ————— **0s** 26ms/step - loss: 3.0330  
Epoch 3/50  
3/3 ————— **0s** 18ms/step - loss: 3.0179  
Epoch 4/50  
3/3 ————— **0s** 20ms/step - loss: 3.0033  
Epoch 5/50  
3/3 ————— **0s** 25ms/step - loss: 2.9773  
Epoch 6/50  
3/3 ————— **0s** 35ms/step - loss: 2.9452  
Epoch 7/50  
3/3 ————— **0s** 34ms/step - loss: 2.8873  
Epoch 8/50  
3/3 ————— **0s** 65ms/step - loss: 2.8286  
Epoch 9/50  
3/3 ————— **0s** 36ms/step - loss: 2.7777  
Epoch 10/50  
3/3 ————— **0s** 29ms/step - loss: 2.8105  
Epoch 11/50  
3/3 ————— **0s** 20ms/step - loss: 2.7791  
Epoch 12/50  
3/3 ————— **0s** 20ms/step - loss: 2.7369  
Epoch 13/50  
3/3 ————— **0s** 64ms/step - loss: 2.7485  
Epoch 14/50  
3/3 ————— **0s** 25ms/step - loss: 2.7649  
Epoch 15/50  
3/3 ————— **0s** 24ms/step - loss: 2.7455  
Epoch 16/50  
3/3 ————— **0s** 31ms/step - loss: 2.6912  
Epoch 17/50  
3/3 ————— **0s** 24ms/step - loss: 2.7197  
Epoch 18/50  
3/3 ————— **0s** 25ms/step - loss: 2.7014  
Epoch 19/50  
3/3 ————— **0s** 20ms/step - loss: 2.6688  
Epoch 20/50  
3/3 ————— **0s** 36ms/step - loss: 2.6845  
Epoch 21/50  
3/3 ————— **0s** 20ms/step - loss: 2.6241  
Epoch 22/50  
3/3 ————— **0s** 21ms/step - loss: 2.6552  
Epoch 23/50  
3/3 ————— **0s** 37ms/step - loss: 2.5843  
Epoch 24/50  
3/3 ————— **0s** 28ms/step - loss: 2.6730  
Epoch 25/50  
3/3 ————— **0s** 35ms/step - loss: 2.5956  
Epoch 26/50  
3/3 ————— **0s** 18ms/step - loss: 2.5933  
Epoch 27/50  
3/3 ————— **0s** 12ms/step - loss: 2.5536

Epoch 28/50  
3/3 ————— 0s 13ms/step - loss: 2.4900  
Epoch 29/50  
3/3 ————— 0s 14ms/step - loss: 2.5217  
Epoch 30/50  
3/3 ————— 0s 12ms/step - loss: 2.4694  
Epoch 31/50  
3/3 ————— 0s 14ms/step - loss: 2.4195  
Epoch 32/50  
3/3 ————— 0s 13ms/step - loss: 2.4173  
Epoch 33/50  
3/3 ————— 0s 14ms/step - loss: 2.4048  
Epoch 34/50  
3/3 ————— 0s 12ms/step - loss: 2.3628  
Epoch 35/50  
3/3 ————— 0s 12ms/step - loss: 2.3708  
Epoch 36/50  
3/3 ————— 0s 12ms/step - loss: 2.3163  
Epoch 37/50  
3/3 ————— 0s 12ms/step - loss: 2.2657  
Epoch 38/50  
3/3 ————— 0s 12ms/step - loss: 2.2793  
Epoch 39/50  
3/3 ————— 0s 12ms/step - loss: 2.2479  
Epoch 40/50  
3/3 ————— 0s 14ms/step - loss: 2.1738  
Epoch 41/50  
3/3 ————— 0s 12ms/step - loss: 2.1553  
Epoch 42/50  
3/3 ————— 0s 12ms/step - loss: 2.1356  
Epoch 43/50  
3/3 ————— 0s 16ms/step - loss: 2.1616  
Epoch 44/50  
3/3 ————— 0s 12ms/step - loss: 2.1351  
Epoch 45/50  
3/3 ————— 0s 14ms/step - loss: 2.0126  
Epoch 46/50  
3/3 ————— 0s 14ms/step - loss: 1.9456  
Epoch 47/50  
3/3 ————— 0s 12ms/step - loss: 1.9893  
Epoch 48/50  
3/3 ————— 0s 12ms/step - loss: 1.9486  
Epoch 49/50  
3/3 ————— 0s 12ms/step - loss: 1.8525  
Epoch 50/50  
3/3 ————— 0s 13ms/step - loss: 1.8525  
Generated text:  
lstm is a yye preeerer nnnenaalllggggpppeessssiiiiintnntuaa

## 11) SENTIMENT ANALYSIS with LSTM:

```
import tensorflow as tf
from tensorflow.keras.datasets import imdb
```

---

```

from tensorflow.keras.preprocessing import sequence
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense

# Load IMDB dataset (only top 10,000 words)
max_features = 10000
maxlen = 200 # cut texts after 200 words
batch_size = 32

print("Loading data...")
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)
print(len(x_train), "train sequences")
print(len(x_test), "test sequences")

# Pad sequences to equal length
x_train = sequence.pad_sequences(x_train, maxlen=maxlen)
x_test = sequence.pad_sequences(x_test, maxlen=maxlen)

# Build LSTM model
model = Sequential()
model.add(Embedding(max_features, 128))
model.add(LSTM(128, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(1, activation="sigmoid"))

# Compile & Train
model.compile(loss="binary_crossentropy", optimizer="adam", metrics=["accuracy"])
print("Training...")
model.fit(x_train, y_train, batch_size=batch_size, epochs=2, validation_data=(x_test, y_test))

# Evaluate
score, acc = model.evaluate(x_test, y_test, batch_size=batch_size)
print("Test score:", score)
print("Test accuracy:", acc)

```

## OUTPUT:

Loading data... 25000 train sequences 25000 test sequences Training... Epoch 1/2 **782/782**  
**558s** 709ms/step - accuracy: 0.7112 - loss: 0.5426 -  
 val\_accuracy: 0.8265 - val\_loss: 0.3970 Epoch 2/2 **782/782**  
**555s** 710ms/step - accuracy: 0.8339 - loss: 0.3815 - val\_accuracy: 0.8284 - val\_loss: 0.3968 **782/782**  
**92s** 118ms/step - accuracy: 0.8264 - loss: 0.4007

Test score: 0.39676380157470703

Test accuracy: 0.8284000158309937

